

A Practical Framework for Better Hotel Comp Decisions

How to balance guest recovery, consistency, and cost before changing policy

Grant McCurdy

Evidence boundary. This prototype uses synthetic hotel operating records and policy outcomes. Official public Santa Monica Proper sources inform the property-relevant recovery menu and guest-facing values only. The work does not use or claim access to Proper Hotels guest records, comp history, internal costs, policies, inventory, or performance.

The decision is not whether to comp less

In luxury hospitality, a comp is both a cost and a service-recovery tool. Cost-only rules can damage guest relationships; unstructured discretion can create inconsistency and unnecessary room-rate erosion. The business question is:

How can managers choose the right gesture for the guest and the situation, while making the cost and reasoning visible?

Recommendation: do not change comp policy from synthetic results. First map the hotel's actual data and decision rules, then run an invisible shadow evaluation for four weeks or 50 eligible recovery cases, whichever is later. Compare prototype recommendations with manager judgment, actual cost, operational feasibility, and observed guest outcomes before deciding whether a controlled policy test is warranted.

The four-week/50-case threshold is an operational discovery target, not a statistically powered sample for proving impact.

Why this sequence is safer

1. **Guest recovery first.** A recommendation must meet proposed recovery, escalation, data-quality, and feasibility rules before cost enters the comparison.
2. **Cost second.** Among options that clear those safeguards, compare actual marginal cost rather than guest-facing price alone.
3. **Real outcomes decide.** Satisfaction recovery, reviews, repeat stays, manager overrides, and contribution margin determine whether the rule deserves a controlled test.

What the current prototype demonstrates

The prototype reconciles synthetic PMS, CRM, service-ticket, comp-ledger, POS, survey, and operating data into recovery cases, then tests five policies under the same proposed rules. The screen checks recovery paths, high-risk escalation, operational feasibility, data holds, and highest-severity review. Uncertain cost and case assumptions are varied repeatedly; every rule must still pass in at least 80% of those runs.

Under these synthetic assumptions, three policies pass the proposed screen. The lowest modeled median among them is approximately \$71 per case. That result shows how a safeguards-first comparison works; it does **not** identify an optimal policy for Proper Hotels.

One case, end to end

A synthetic returning guest experiences a severity-4 room-readiness delay for which hotel responsibility is modeled at 93%. The engine recommends a **\$180 Calabria or Palma dining credit**, with an assumed internal-cost range of **\$45-\$108**.

Cost is compared only after recovery safeguards

Modeled internal cost per synthetic recovery case




Figure 1: Only the three policies that pass every proposed synthetic screening rule are shown. The dot is the middle modeled cost per synthetic case; the line contains the central 90% of results after uncertain assumptions are varied 5,000 times. This is a sensitivity check, not observed Proper Hotels costs, savings, or a probability of business success. Source: versioned synthetic policy mart, 430 cases per policy.

The explanation is visible: the issue is hotel-responsible, high severity, and recoverable before checkout; the dining gesture has high perceived value at a lower assumed cost. Without the operating-availability signal, the model would choose a \$360 room upgrade; without the property-fit signal, it would choose \$100 late checkout. Managers can challenge that reasoning rather than accept a black-box score.

A controlled path to a real decision

Stage	Business question	Decision gate
1. Data and policy workshop	Can failures, comps, costs, constraints, and outcomes be joined reliably?	Proceed only after definitions, ownership, and material data gaps are agreed.
2. Invisible shadow evaluation	Are recommendations guest-appropriate, feasible, explainable, and economically credible beside manager decisions?	Revise or stop when safety, data-quality, feasibility, or trust checks fail.
3. Controlled-test design	Is there enough observed evidence to test manager-assisted use?	Pre-register outcomes and sample size; do not adopt policy from shadow agreement alone.

The shadow record should capture the failure and timing, guest context, manager decision and override, gesture availability, actual marginal cost, and post-recovery outcomes. Those data distinguish intelligent generosity from under-recovery and avoidable leakage.

Proposed next conversation

Use a 90-minute working session to map the current comp workflow, define what “good recovery” means, identify the relevant systems and owners, and select a small set of shadow metrics. The immediate output is a shared data map and validation plan. A manager-facing tool or policy change comes only after the real evidence supports it.